

# Memory Allocators

Asim Maharjan  
Madhav Aryal


# Revision

- Virtual Memory
- Pointer Arithmetic
- Bitwise Operations

# Contents

- Allocating Virtual Memory (Windows and Linux)
- Memory Alignment
- Types of Allocators
  - Bump Allocator
  - Stack Allocator
  - Free list Allocator

# Virtual Memory



# Allocating/Obtaining Virtual Memory

- Memory is a resource managed by the Operating System
- Use system calls to request OS for more memory
- The OS will allocate pages, creates mappings and returns back a pointer to memory
- Need to utilize OS specific functions
  - *mmap* for Linux based systems
  - *VirtualAlloc* for Windows systems

# Page Allocation

- OS will always allocated virtual memory in terms of pages.
- Consequence: Memory size allocated will always be a multiple of the Page size of the operating system
- Example, if page size of OS is 4kB, and we request a page having 2kB, then OS will return page of size 4kB regardless.

# Memory Protections

- Provided by the OS.
- Three common types of protections:
  - Read Protection
  - Write Protection
  - Execute Protection


# The mmap system call

- Used in Linux based system
- Maps files or devices into memory i.e, creates a memory mapping
- Enables memory mapped I/O



# Memory Mapped I/O

- I/O in which devices/files are accessed through memory addresses (pointers)
- Common example is file I/O which utilizes your HDD or SSD
- Any changes made to the memory through the pointer will also be reflected into the file.



# Relation with Virtual Memory

- Virtual pages are stored in disk
- *mmap* allows to create mappings to files on disk
- Creating mappings to virtual pages
- Named as **anonymous mapping**

# Using mmap for Memory Allocation

```
void *mmap( void *addr, size_t length, int prot, int flags, int fd,  
            off_t offset);
```

Returns a pointer to memory on success, MAP\_FAILED on failure.

addr	Starting address of memory to allocate. If set to NULL, OS will return a suitable pointer instead.
length	Size of the memory in bytes we want to allocate.
prot	Protection flags
flags	Controls properties of the mapping
fd	Set to -1 for our purposes
offset	Set to 0 for our purposes

# mmap: An Example

addr set to null

size of  
4kB

Enable Read and  
write to  
allocated memory

```
size_t size = 4096;
void *mem = mmap( NULL, 4096, PROT_READ|PROT_WRITE,
                  MAP_ANON|MAP_PRIVATE, -1, 0 );

if ( mem == MAP_FAILED ){
    // .. error
    exit(1);
}
```

Create a  
private and  
anonymous map

fd and offset  
set to -1 and 0

# The munmap system call

- Undos a mapping created by mmap
- Basically, deallocates the pages allocated with mmap

```
int munmap( void *addr, size_t length );
```

- Undos the mapping of all pages that are in address range of **addr** to **addr+length**
- Returns 0 on success and -1 on error.

## mmap Full Example

```
size_t size = 4096;
void *mem = mmap( NULL, size, PROT_READ|PROT_WRITE,
                  MAP_ANON|MAP_PRIVATE, -1, 0 );

if ( mem == MAP_FAILED ){
    // .. error
    exit(1);
}

if ( munmap( mem, size ) == -1 ){
    // error again
}
```

# The VirtualAlloc function

- Windows counterpart of mmap
- “Reserve, commit or change the state of the region of memory within the virtual address space of a specified process.”

# States of Memory

- Memory in Windows can be in 3 states: Reserved, Committed and Free
- Reserved
  - Only the address range is reserved and thus cannot be used by others
  - No pages are allocated
- Committed
  - Pages are allocated
- Free
  - The default state
  - Available for allocation



# The VirtualAlloc system call

```
LPVOID VirtualAlloc(  
    [in, optional] LPVOID lpAddress,  
    [in]           SIZE_T dwSize,  
    [in]           DWORD  flAllocationType,  
    [in]           DWORD  flProtect  
);
```

- Needs header <Memoryapi.h>
- Returns an LPVOID type return value. LPVOID means Long Pointer to Void, i.e, void \*
- Returns NULL when fails otherwise returns starting address of allocated memory

## Side note on Windows Type Naming

LPVOID	Long Pointer to VOID type, i.e, a void *
SIZE_T	Equivalent to size_t
DWORD	A 32-bit unsigned integer
[in]	The parameter acts as an input to the function

# The VirtualAlloc system call

```
LPVOID VirtualAlloc(  
    [in, optional] LPVOID lpAddress,  
    [in]           SIZE_T dwSize,  
    [in]           DWORD  flAllocationType,  
    [in]           DWORD  flProtect  
);
```

lpAddress	Similar to that of mmap
dwSize	Size of memory to be allocated
flAllocationType	Set to MEM_RESERVE MEM_COMMIT
flProtect	Protection Flags

# VirtualAlloc: An Example

```
SIZE_T size = 4096;
void *mem = VirtualAlloc(NULL, size,
                        MEM_COMMIT|MEM_RESERVE, PAGE_READWRITE );

if ( mem == NULL ){
    //.... error
    exit(0);
}
// Success: Use mem
// .....
```


# VirtualFree

- Decommits/Deallocates memory previously allocated by VirtualAlloc

```
BOOL VirtualFree(  
    [in] LPVOID lpAddress,  
    [in] SIZE_T dwSize,  
    [in] DWORD dwFreeType  
);
```

lpAddress	Starting address of memory to be freed
dwSize	Size of memory to be freed
dwFreeType	MEM_DECOMMIT or MEM_RELEASE

# States of Memory



# MEM\_DECOMMIT and MEM\_RELEASE

- MEM\_DECOMMIT sends pages back to MEM\_RESERVED state
- MEM\_RELEASE frees actual pages
  - dwSize **must** be 0 when MEM\_RELEASE is used.

## Virtual Free: An example

```
BOOL ret = VirtualFree(mem, 0, MEM_RELEASE);  
if ( !ret ){  
    //... handle error  
}  
//.. No error
```

```
BOOL ret = VirtualFree(mem, 4096, MEM_DECOMMIT);  
if ( !ret ){  
    //... handle error  
}  
  
ret = VirtualFree( mem, 0, MEM_RELEASE );  
//.. No error
```



# Memory Alignment

- A variable of a certain data type stored in memory follows an “alignment rule”
- The rule states which addresses are considered “valid” for a certain data type
- Some examples:
  - A 32-bit integer is always stored at an address that is a multiple of 4
  - A 64-bit integer is always stored at an address that is a multiple of 8
  - and so on..
- Dependent on the size of the data type.
- Access from an unaligned memory leads to slower code

# Aligning to a Power of Two

```
uint64_t align_up( uint64_t size, uint64_t align ){  
    return ( size + ( align - 1 ) ) & ~(align-1);  
}
```

- The above functions “aligns” the value of size to the power of 2 given by align
- `align_up( 23, 8 ) = 24`
- `align_up(27,8) = 32`
- `align_up(16,8) = 16`

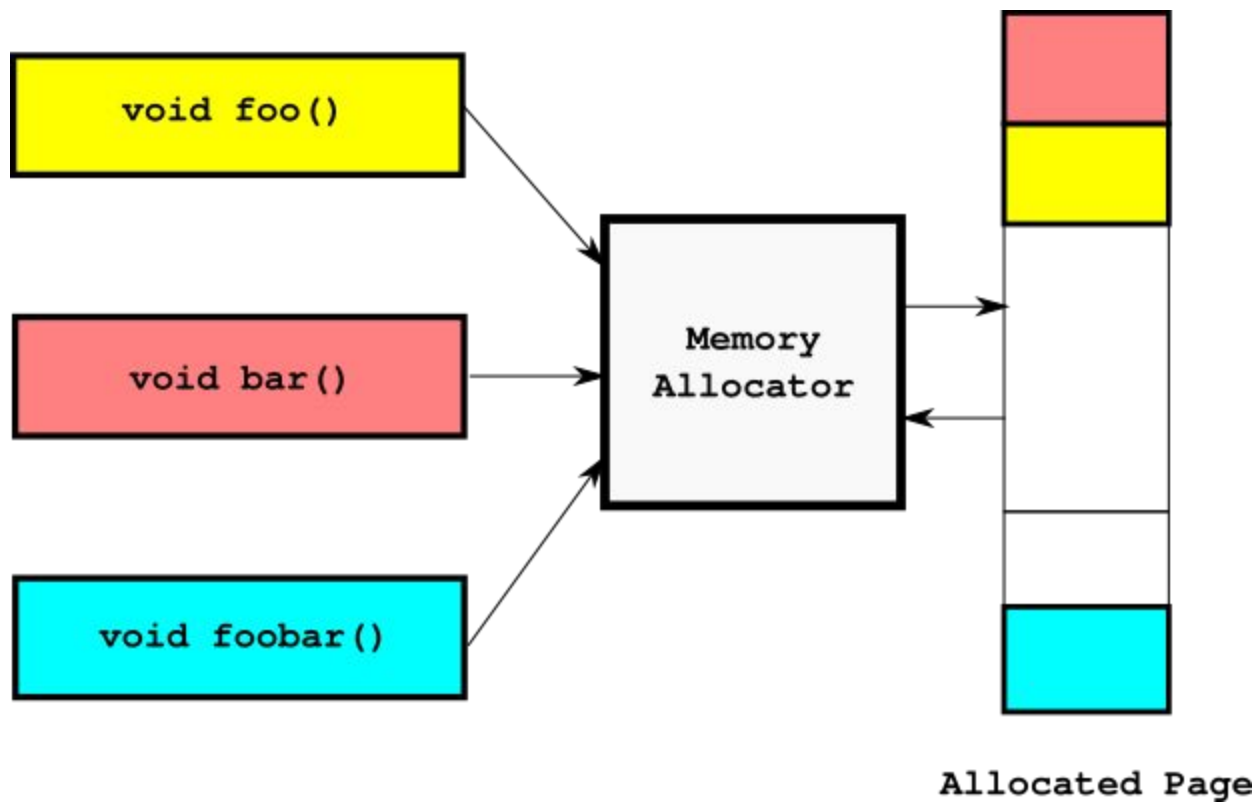
# Aligning Up to a power of Two: Breakdown

- What does  $X \& \sim(Y-1)$  do when  $Y$  is a power of two?
  - Hint: Write  $Y-1$  in binary
- What is the value of  $(X + (\text{align}-1))$ ?
  - Check for different values of  $X$ , when  $X$  is a multiple of  $\text{align}$  and when it is not

**BREAK**

# Memory Allocators

- After allocating virtual memory from the OS, we are left with a single big chunk of memory
- We still need some way to distribute and manage this chunk of memory among different parts of our program
- Main functionalities:
  - An 'alloc' function to allow our program to request memory
  - A 'free' function to return memory back so that it can be reused.
  - Allows for addresses of arbitrary alignment



# Memory Allocators: Types


- We discuss 4 different types of allocators
- In Increasing order of their complexity and flexibility:
  - Bump Allocator
  - Stack Allocator
  - Pooled style Free list Allocator
  - Generic Free List Allocator
- All allocators have their own way of performing memory allocation and deallocation

# Bump Allocators

- Simplest kind of memory allocator
- Allows for memory allocations
- Allows limited freeing of memory
- It is **fast** and can be reused easily.



# Bump Allocator: The Bigger Picture




- Different color indicate memory allocated by different parts of the program
- Gray region has address of 0x100, Red(maroon) has address of 0x200 and so on.
- The next allocation made will have the memory address 0x290.
- Alignment is not considered (yet).

# Bump Allocator: Data Structure

- **Current** pointer, which points to base of unallocated memory
- Pointer to the memory allocated from the OS.
- Size of the memory allocated from the OS.

# Bump Allocator: Data Structure

- In our previous example:
  - **Current** = 0x290
  - Pointer to memory allocated from OS = 0x100
  - Size of allocated memory = 1024 bytes




# Bump Allocator: Allocation

- Consider allocation of X bytes of memory with a given alignment
- The steps involved include:
  - Aligning value of **current** pointer to the given alignment
  - Saving value of **current** pointer
  - Incrementing value of **current** pointer by X bytes
  - Performing bounds check
  - Returning the allocated memory (i.e, the save **current** pointer)
- Consider the allocation of 37 bytes initially, with an alignment of 8 bytes


# Bump Allocator: Allocation

- Initial State:
  - **Current** = 0x100
  - Memory pointer = 0x100
  - Size = 1024 bytes
- Notice that **current** pointer is already aligned to 8 bytes




# Bump Allocator: Allocation

- The **current** pointer is incremented by 37 bytes.
- Bounds checking the **current** pointer
  - Lies inside the allocated memory from the OS
  - Ensures that 37 bytes can be utilized
- The pointer value of 0x100 is returned




## Bump Allocator: Second Allocation

- Now consider allocation of another 100 bytes with 8 byte alignment
- Align **current** pointer to 8 bytes
  - $\text{align}(0x125, 8) = 0x128$



# Bump Allocator: Second Allocation

- Increment **current** pointer by 100 bytes
- Bounds checking the **current** pointer
  - Lies inside the allocated memory from the OS
- The pointer value of 0x128 is returned

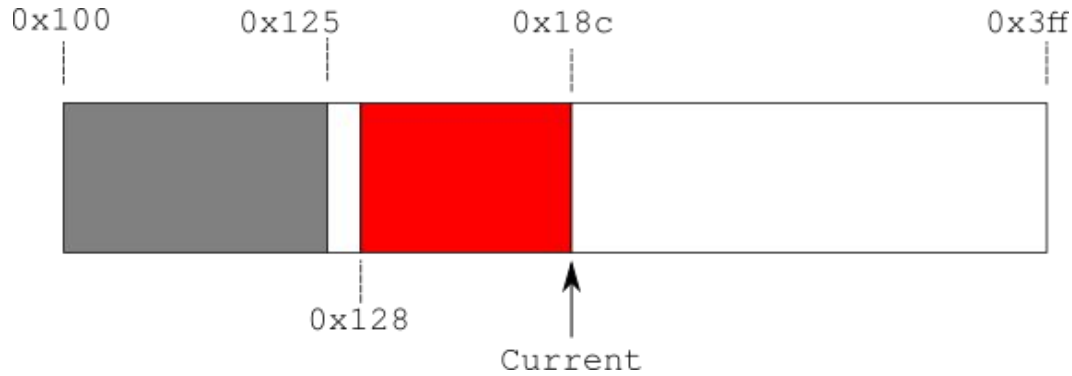





# Bump Allocator: Freeing Memory

- Limited when it comes to freeing memory
- Can only free memory from the top
- Size of memory to be freed needs to be passed
- Freeing is simple
  - Simply decrement the **current** pointer by the size to be freed
  - Bounds checking still needs to be done

# Bump Allocator: Freeing Memory Example



After Freeing by 100 bytes:



# Bump Allocator: Resetting

- Instead of freeing from the top, a more common operation is resetting of the bump allocator
- Simply sets the allocator to its initial state
- Easily Re-use the allocator for other purposes


# Bump Allocator: Summary

- Advantages:
  - Simple and Fast
- Drawbacks
  - Cannot grow in size
  - Limited freeing
- Useful for temporary memory, i.e, applications where we need memory only for a relatively small number of operations.

# Stack Allocators

- Similar to a Bump Allocator in their operation
- However, it keeps track of every allocation
- Freeing however, still can only be done from the top


# Stack Allocators: The Bigger Picture



- Allocations are tracked using a linked list of pointers that goes backwards


# Stack Allocator: Data Structure

- **Current** pointer, which points to base of unallocated memory
- Pointer to the memory allocated from the OS.
- Size of the memory allocated from the OS.
- **Top** pointer, which points to the beginning of the linked list structure



# Stack Allocator: Allocation


1. Align the **current** point suitably so that we can store a pointer.






# Stack Allocator: Allocation

2. Store the current **Top** value in that address




# Stack Allocator: Allocation

3. Set the **Top** value to the **Current** pointer




# Stack Allocator: Allocation

4. Point to the next memory address after. This is the address to be returned back. Any required alignment should be done in this step.



# Stack Allocator: Allocation

5. Allocate the required number of bytes as done in bump allocator.




# Stack Allocator: Freeing Memory

- Stack allocator provides a **Pop()** function that can be used to free memory at the top
- The **Pop()** does not need a size argument to be passed


# Stack Allocator: Freeing Memory

1. Set **Current** = **Top**



# Stack Allocator: Freeing Memory

## 2. Set **Top** = \***Top**



# Stack Allocator: Summary

- Not very useful compared to bump allocator
- Added a lot complexity for addition of a very small functionality
- Total size of allocator is fixed.



# Free List Allocator

- Uses a linked list to store memory chunks
- Memory chunks are of fixed size
- Allocations of only a fixed size can occur using this kind of free list
- Allows for general freeing of memory blocks
- We do NOT consider the case for an arbitrary aligned memory.

# Free List Allocator: Algorithm


1. Start with an empty page

0x200




# Free List Allocator: Algorithm

2. Create a linked list of memory chunks. Here, we have assumed a chunk size of 256 (0x100) bytes.




# Free List Allocator: Allocation

3. To allocate a chunk, remove the Top node from the linked list and return it.




# Free List Allocator: Allocation

Here's another example:



# Free List Allocator: Freeing Memory

4. To free memory, we simply insert it into the HEAD of the linked list and store the pointer to the next chunk.



# Free List Allocator: Getting New Memory

- What happens when free list runs out of chunks?
  - Simply allocate more Virtual Memory and perform the same steps as above
- Need to keep track of all the virtual memory allocated so far
  - Required to free all the created mappings

# Free List Allocator: Summary

- Allows for freeing of individual allocations
- Can grow in size
- Drawbacks
  - Allocation size is fixed